

NAME

ovn-northd – Open Virtual Network central control daemon

SYNOPSIS

ovn-northd [*options*]

DESCRIPTION

ovn-northd is a centralized daemon responsible for translating the high-level OVN configuration into logical configuration consumable by daemons such as **ovn-controller**. It translates the logical network configuration in terms of conventional network concepts, taken from the OVN Northbound Database (see **ovn-nb**(5)), into logical datapath flows in the OVN Southbound Database (see **ovn-sb**(5)) below it.

CONFIGURATION

ovn-northd requires a connection to the Northbound and Southbound databases. The default is **db.sock** in the local Open vSwitch's "run" directory. This may be overridden with the following commands:

- **--ovnnb-db=database**
The database containing the OVN Northbound Database.
- **--ovsdb-db=database**
The database containing the OVN Southbound Database.

The *database* argument must take one of the following forms:

- **ssl:ip:port**
The specified SSL *port* on the host at the given *ip*, which must be expressed as an IP address (not a DNS name) in IPv4 or IPv6 address format. If *ip* is an IPv6 address, then wrap *ip* with square brackets, e.g.: **ssl:::1:6640**. The **--private-key**, **--certificate**, and **--ca-cert** options are mandatory when this form is used.
- **tcp:ip:port**
Connect to the given TCP *port* on *ip*, where *ip* can be IPv4 or IPv6 address. If *ip* is an IPv6 address, then wrap *ip* with square brackets, e.g.: **tcp:::1:6640**.
- **unix:file**
On POSIX, connect to the Unix domain server socket named *file*.
On Windows, connect to a localhost TCP port whose value is written in *file*.

RUNTIME MANAGEMENT COMMANDS

ovs-appctl can send commands to a running **ovn-northd** process. The currently supported commands are described below.

- exit** Causes **ovn-northd** to gracefully terminate.

LOGICAL FLOW TABLE STRUCTURE

One of the main purposes of **ovn-northd** is to populate the **Logical_Flow** table in the **OVN_Southbound** database. This section describes how **ovn-northd** does this for switch and router logical datapaths.

Logical Switch Datapaths

Ingress Table 0: Admission Control and Ingress Port Security

Ingress table 0 contains these logical flows:

- Priority 100 flows to drop packets with VLAN tags or multicast Ethernet source addresses.
- Priority 50 flows that implement ingress port security for each enabled logical port. For logical ports on which port security is enabled, these match the **inport** and the valid **eth.src** address(es) and advance only those packets to the next flow table. For logical ports on which port security is not enabled, these advance all packets that match the **inport**.

There are no flows for disabled logical ports because the default-drop behavior of logical flow tables causes packets that ingress from them to be dropped.

Ingress Table 1: from-lport Pre-ACLs

Ingress table 1 prepares flows for possible stateful ACL processing in table 2. It contains a priority-0 flow that simply moves traffic to table 2. If stateful ACLs are used in the logical datapath, a priority-100 flow is added that sends IP packets to the connection tracker before advancing to table 2.

Ingress table 2: from-lport ACLs

Logical flows in this table closely reproduce those in the **ACL** table in the **OVN_Northbound** database for the **from-lport** direction. **allow** ACLs translate into logical flows with the **next;** action, **allow-related** ACLs translate into logical flows with the **ct_next;** action, other ACLs translate to **drop;**. The **priority** values from the **ACL** table are used directly.

Ingress table 2 also contains a priority 0 flow with action **next;**, so that ACLs allow packets by default. If the logical datapath has a stateful ACL, the following flows will also be added:

- A priority-1 flow to commit IP traffic to the connection tracker. This is needed for the default allow policy because, while the initiator's direction may not have any stateful rules, the server's may and then its return traffic would not be known and marked as invalid.
- A priority-65535 flow that allows any traffic that has been committed to the connection tracker (i.e., established flows).
- A priority-65535 flow that allows any traffic that is considered related to a committed flow in the connection tracker (e.g., an ICMP Port Unreachable from a non-listening UDP port).
- A priority-65535 flow that drops all traffic marked by the connection tracker as invalid.

Ingress Table 3: ARP responder

This table implements ARP responder for known IPs. It contains these logical flows:

- Priority-100 flows to skip ARP responder if inport is of type **localnet**, and advances directly to table 3.
- Priority-50 flows that matches ARP requests to each known IP address *A* of logical port *P*, and respond with ARP replies directly with corresponding Ethernet address *E*:

```
eth.dst = eth.src;
eth.src = E;
arp.op = 2; /* ARP reply. */
arp.tha = arp.sha;
arp.sha = E;
arp.tpa = arp.spa;
arp.spa = A;
output = P;
inport = ""; /* Allow sending out inport. */
output;
```

These flows are omitted for logical ports (other than router ports) that are down.

- One priority-0 fallback flow that matches all packets and advances to table 4.

Ingress Table 4: Destination Lookup

This table implements switching behavior. It contains these logical flows:

- A priority-100 flow that outputs all packets with an Ethernet broadcast or multicast **eth.dst** to the **MC_FLOOD** multicast group, which **ovn-northd** populates with all enabled logical ports.

- One priority-50 flow that matches each known Ethernet address against **eth.dst** and outputs the packet to the single associated output port.
- One priority-0 fallback flow that matches all packets and outputs them to the **MC_UNKNOWN** multicast group, which **ovn-northd** populates with all enabled logical ports that accept unknown destination packets. As a small optimization, if no logical ports accept unknown destination packets, **ovn-northd** omits this multicast group and logical flow.

Egress Table 0: to-lport Pre-ACLs

This is similar to ingress table 1 except for **to-lport** traffic.

Egress Table 1: to-lport ACLs

This is similar to ingress table 2 except for **to-lport** ACLs.

Egress Table 2: Egress Port Security

This is similar to the ingress port security logic in ingress table 0, but with important differences. Most obviously, **outport** and **eth.dst** are checked instead of **inport** and **eth.src**. Second, packets directed to broadcast or multicast **eth.dst** are always accepted instead of being subject to the port security rules; this is implemented through a priority-100 flow that matches on **eth.mcast** with action **output;**. Finally, to ensure that even broadcast and multicast packets are not delivered to disabled logical ports, a priority-150 flow for each disabled logical **outport** overrides the priority-100 flow with a **drop;** action.

Logical Router Datapaths

Ingress Table 0: L2 Admission Control

This table drops packets that the router shouldn't see at all based on their Ethernet headers. It contains the following flows:

- Priority-100 flows to drop packets with VLAN tags or multicast Ethernet source addresses.
- For each enabled router port *P* with Ethernet address *E*, a priority-50 flow that matches **inport == P && (eth.mcast || eth.dst == E)**, with action **next;**.

Other packets are implicitly dropped.

Ingress Table 1: IP Input

This table is the core of the logical router datapath functionality. It contains the following flows to implement very basic IP host functionality.

- L3 admission control: A priority-100 flow drops packets that match any of the following:
 - **ip4.src[28..31] == 0xe** (multicast source)
 - **ip4.src == 255.255.255.255** (broadcast source)
 - **ip4.src == 127.0.0.0/8 || ip4.dst == 127.0.0.0/8** (localhost source or destination)
 - **ip4.src == 0.0.0.0/8 || ip4.dst == 0.0.0.0/8** (zero network source or destination)
 - **ip4.src** is any IP address owned by the router.
 - **ip4.src** is the broadcast address of any IP network known to the router.
- ICMP echo reply. These flows reply to ICMP echo requests received for the router's IP address. Let *A* be an IP address or broadcast address owned by a router port. Then, for each *A*, a priority-90 flow matches on **ip4.dst == A** and **icmp4.type == 8 && icmp4.code == 0** (ICMP echo request). These flows use the following actions where, if *A* is unicast, then *S* is *A*, and if *A* is broadcast, *S* is the router's IP address in *A*'s network:

```
ip4.dst = ip4.src;
ip4.src = S;
ip.ttl = 255;
```

```
icmp4.type = 0;
inport = ""; /* Allow sending out inport. */
next;
```

Similar flows match on **ip4.dst == 255.255.255.255** and each individual **inport**, and use the same actions in which *S* is a function of **inport**.

- ARP reply. These flows reply to ARP requests for the router's own IP address. For each router port *P* that owns IP address *A* and Ethernet address *E*, a priority-90 flow matches **inport == P && arp.tpa == A && arp.op == 1** (ARP request) with the following actions:

```
eth.dst = eth.src;
eth.src = E;
arp.op = 2; /* ARP reply. */
arp.tha = arp.sha;
arp.sha = E;
arp.tpa = arp.spa;
arp.spa = A;
output = P;
inport = ""; /* Allow sending out inport. */
output;
```

- UDP port unreachable. Priority-80 flows generate ICMP port unreachable messages in reply to UDP datagrams directed to the router's IP address. The logical router doesn't accept any UDP traffic so it always generates such a reply.

These flows should not match IP fragments with nonzero offset.

Details TBD. Not yet implemented.

- TCP reset. Priority-80 flows generate TCP reset messages in reply to TCP datagrams directed to the router's IP address. The logical router doesn't accept any TCP traffic so it always generates such a reply.

These flows should not match IP fragments with nonzero offset.

Details TBD. Not yet implemented.

- Protocol unreachable. Priority-70 flows generate ICMP protocol unreachable messages in reply to packets directed to the router's IP address on IP protocols other than UDP, TCP, and ICMP.

These flows should not match IP fragments with nonzero offset.

Details TBD. Not yet implemented.

- Drop other IP traffic to this router. These flows drop any other traffic destined to an IP address of this router that is not already handled by one of the flows above, which amounts to ICMP (other than echo requests) and fragments with nonzero offsets. For each IP address *A* owned by the router, a priority-60 flow matches **ip4.dst == A** and drops the traffic.

The flows above handle all of the traffic that might be directed to the router itself. The following flows (with lower priorities) handle the remaining traffic, potentially for forwarding:

- Drop Ethernet local broadcast. A priority-50 flow with match **eth.bcast** drops traffic destined to the local Ethernet broadcast address. By definition this traffic should not be forwarded.
- Drop IP multicast. A priority-50 flow with match **ip4.mcast** drops IP multicast traffic.
- ICMP time exceeded. For each router port *P*, whose IP address is *A*, a priority-40 flow with match **inport == P && ip.ttl == {0, 1} && !ip.later_frag** matches packets whose TTL has expired, with the following actions to send an ICMP time exceeded reply:

```
icmp4 {
    icmp4.type = 11; /* Time exceeded. */
    icmp4.code = 0; /* TTL exceeded in transit. */
    ip4.dst = ip4.src;
    ip4.src = A;
    ip.ttl = 255;
    next;
};
```

Not yet implemented.

- TTL discard. A priority-30 flow with match **ip.ttl == {0, 1}** and actions **drop**; drops other packets whose TTL has expired, that should not receive a ICMP error reply (i.e. fragments with nonzero offset).
- Next table. A priority-0 flows match all packets that aren't already handled and uses actions **next**; to feed them to the ingress table for routing.

Ingress Table 2: IP Routing

A packet that arrives at this table is an IP packet that should be routed to the address in **ip4.dst**. This table implements IP routing, setting **reg0** to the next-hop IP address (leaving **ip4.dst**, the packet's final destination, unchanged) and advances to the next table for ARP resolution.

This table contains the following logical flows:

- Routing table. For each route to IPv4 network N with netmask M , a logical flow with match **ip4.dst == N/M** , whose priority is the number of 1-bits in M , has the following actions:

```
ip.ttl--;
reg0 = G;
next;
```

(Ingress table 1 already verified that **ip.ttl--**; will not yield a TTL exceeded error.)

If the route has a gateway, G is the gateway IP address, otherwise it is **ip4.dst**.

- Destination unreachable. For each router port P , which owns IP address A , a priority-0 logical flow with match **in_port == P && !ip.later_frag && !icmp** has the following actions:

```
icmp4 {
    icmp4.type = 3; /* Destination unreachable. */
    icmp4.code = 0; /* Network unreachable. */
    ip4.dst = ip4.src;
    ip4.src = A;
    ip.ttl = 255;
    next(2);
};
```

(The **!icmp** check prevents recursion if the destination unreachable message itself cannot be routed.)

These flows are omitted if the logical router has a default route, that is, a route with netmask 0.0.0.0.

Ingress Table 3: ARP Resolution

Any packet that reaches this table is an IP packet whose next-hop IP address is in **reg0**. (**ip4.dst** is the final destination.) This table resolves the IP address in **reg0** into an output port in **outputport** and an Ethernet address in **eth.dst**, using the following flows:

- Known MAC bindings. For each IP address A whose host is known to have Ethernet address HE and reside on router port P with Ethernet address PE , a priority-200 flow

with match **reg0 == A** has the following actions:

```
eth.src = PE;
eth.dst = HE;
output = P;
output;
```

MAC bindings can be known statically based on data in the **OVN_Northbound** database. For router ports connected to logical switches, MAC bindings can be known statically from the **addresses** column in the **Logical_Port** table. For router ports connected to other logical routers, MAC bindings can be known statically from the **mac** and **network** column in the **Logical_Router_Port** table.

- Unknown MAC bindings. For each non-gateway route to IPv4 network N with netmask M on router port P that owns IP address A and Ethernet address E , a logical flow with match **ip4.dst == N/M**, whose priority is the number of 1-bits in M , has the following actions:

```
arp {
  eth.dst = ff:ff:ff:ff:ff:ff;
  eth.src = E;
  arp.sha = E;
  arp.tha = 00:00:00:00:00:00;
  arp.spa = A;
  arp.tpa = ip4.dst;
  arp.op = 1; /* ARP request. */
  output = P;
  output;
};
```

TBD: How to install MAC bindings when an ARP response comes back. (Implement a "learn" action?)

Not yet implemented.

Egress Table 0: Delivery

Packets that reach this table are ready for delivery. It contains priority-100 logical flows that match packets on each enabled logical router port, with action **output**;